
IMAGE PROCESSING USING RCNN CLASSIFIERS

Table of Contents

AUTHOR	1
Load, Store and Explore the data	1
Divide the dataset into Training, Validation and Testing	1
List the image categories	1
Properties of CIFAR10 dataset	2
CNN Classifier - Middle Layer	2
Final Layer of CNN	3
Combine input, middle and last layers	4
Initialise first convolution layer weight	5
Train CNN using CIFAR-10 Data	5
Validate CIFAR-10 Network training	5

AUTHOR

S Sai Suryateja Vellore Institute of Technology, Vellore

Load, Store and Explore the data

```
cifar10Data = tempdir;  
url = 'https://www.cs.toronto.edu/~kriz/cifar-10-matlab.tar.gz';  
helperCIFAR10Data.download(url, cifar10Data);
```

Divide the dataset into Training, Validation and Testing

```
[trainingImages, trainingLabels, testImages, testLabels] =  
helperCIFAR10Data.load(cifar10Data);
```

List the image categories

```
numImageCategories = 10;  
categories(trainingLabels)
```

```
ans =
```

```
10x1 cell array
```

```
{ 'airplane' }  
{ 'automobile' }  
{ 'bird' }  
{ 'cat' }  
{ 'deer' }  
{ 'dog' }  
{ 'frog' }  
{ 'horse' }  
{ 'ship' }  
{ 'truck' }
```

Properties of CIFAR10 dataset

Create the image input layer for 32x32x3 CIFAR-10 images.

```
[height,width,numChannels, ~] = size(trainingImages);  
  
imageSize = [height width numChannels];  
inputLayer = imageInputLayer(imageSize)
```

inputLayer =

ImageInputLayer with properties:

```
    Name: ''  
    InputSize: [32 32 3]
```

Hyperparameters

```
    DataAugmentation: 'none'  
    Normalization: 'zerocenter'  
    AverageImage: []
```

CNN Classifier - Middle Layer

Convolutional layer parameters

```
filterSize = [5 5];  
numFilters = 32;  
middleLayers = [  
% The first convolutional layer has a bank of 32 5x5x3 filters.  
% A symmetric padding of 2 pixels is added to ensure that image  
% borders are included in the processing. This is important to avoid  
% information at the borders being washed away too early in the  
% network.  
convolution2dLayer(filterSize,numFilters,'Padding',2)  
% Note that the third dimension of the filter can be omitted because  
% it is automatically deduced based on the connectivity of the network.  
% In this case because this layer follows the image layer, the third
```

```
dimension must be 3 to match the number of channels in the input
image.
% Next add the ReLU layer:
reluLayer()
% Follow it with a max pooling layer that has a 3x3 spatial pooling
area and a stride of 2 pixels. This down-samples the data dimensions
from 32x32 to 15x15.
maxPooling2dLayer(3,'Stride',2)
% Repeat the 3 core layers to complete the middle of the network.
convolution2dLayer(filterSize,numFilters,'Padding',2)
reluLayer()
maxPooling2dLayer(3, 'Stride',2)
convolution2dLayer(filterSize,2 * numFilters,'Padding',2)
reluLayer()
maxPooling2dLayer(3,'Stride',2)
]
```

middleLayers =

9x1 Layer array with layers:

1	''	Convolution	32 5x5 convolutions with stride [1 1] and padding [2 2 2 2]
2	''	ReLU	ReLU
3	''	Max Pooling	3x3 max pooling with stride [2 2] and padding [0 0 0 0]
4	''	Convolution	32 5x5 convolutions with stride [1 1] and padding [2 2 2 2]
5	''	ReLU	ReLU
6	''	Max Pooling	3x3 max pooling with stride [2 2] and padding [0 0 0 0]
7	''	Convolution	64 5x5 convolutions with stride [1 1] and padding [2 2 2 2]
8	''	ReLU	ReLU
9	''	Max Pooling	3x3 max pooling with stride [2 2] and padding [0 0 0 0]

Final Layer of CNN

```
finalLayers = [
% Add a fully connected layer with 64 output neurons. The output size
of this layer will be an array with a length of 64.
fullyConnectedLayer(64)
% Add an ReLU non-linearity.
reluLayer
% Add the last fully connected layer. At this point, the network must
produce 10 signals that can be used to measure whether the input
image belongs to one category or another. This measurement is made
using the subsequent loss layers.
fullyConnectedLayer(numImageCategories)
% Add the softmax loss layer and classification layer. The final
layers use the output of the fully connected layer to compute the
```

categorical probability distribution over the image classes. During the training process, all the network weights are tuned to minimize the loss over this categorical distribution.

```
softmaxLayer
classificationLayer
]
```

```
finalLayers =
```

```
5x1 Layer array with layers:
```

1	''	Fully Connected	64 fully connected layer
2	''	ReLU	ReLU
3	''	Fully Connected	10 fully connected layer
4	''	Softmax	softmax
5	''	Classification Output	crossentropyex

Combine input, middle and last layers

```
layers = [
    inputLayer
    middleLayers
    finalLayers
]
```

```
layers =
```

```
15x1 Layer array with layers:
```

1	''	Image Input	32x32x3 images with 'zerocenter'
normalization			
2	''	Convolution	32 5x5 convolutions with stride
[1 1] and padding [2 2 2 2]			
3	''	ReLU	ReLU
4	''	Max Pooling	3x3 max pooling with stride [2
2] and padding [0 0 0 0]			
5	''	Convolution	32 5x5 convolutions with stride
[1 1] and padding [2 2 2 2]			
6	''	ReLU	ReLU
7	''	Max Pooling	3x3 max pooling with stride [2
2] and padding [0 0 0 0]			
8	''	Convolution	64 5x5 convolutions with stride
[1 1] and padding [2 2 2 2]			
9	''	ReLU	ReLU
10	''	Max Pooling	3x3 max pooling with stride [2
2] and padding [0 0 0 0]			
11	''	Fully Connected	64 fully connected layer
12	''	ReLU	ReLU
13	''	Fully Connected	10 fully connected layer
14	''	Softmax	softmax
15	''	Classification Output	crossentropyex

Initialise first convolution layer weight

```
layers(2).Weights = 0.0001 * randn([filterSize numChannels  
    numFilters]);
```

Train CNN using CIFAR-10 Data

Set the network training options

```
opts = trainingOptions('sgdm', ...  
    'Momentum', 0.9, ...  
    'InitialLearnRate', 0.001, ...  
    'LearnRateSchedule', 'piecewise', ...  
    'LearnRateDropFactor', 0.1, ...  
    'LearnRateDropPeriod', 8, ...  
    'L2Regularization', 0.004, ...  
    'MaxEpochs', 40, ...  
    'MiniBatchSize', 128, ...  
    'Verbose', true);  
% A trained network is loaded from disk to save time when running the  
% example. Set this flag to true to train the network.  
doTraining = false;  
if doTraining  
    cifar10Net = trainNetwork(trainingImages, trainingLabels, layers,  
        opts);  
else  
    load('rcnnStopSigns.mat', 'cifar10Net')  
end
```


Validate CIFAR-10 Network training

Extract the first convolutional layer weights

```
w = cifar10Net.Layers(2).Weights;  
% rescale the weights to the range [0, 1] for better visualization  
w = rescale(w);  
figure(1)  
montage(w)  
% Run the network on the test set.  
YTest = classify(cifar10Net, testImages);  
% Calculate the accuracy.  
accuracy = sum(YTest == testLabels)/numel(testLabels)
```

accuracy =

0.7456



Published with MATLAB® R2019a